


Memory traffic, not arithmetic, sets the latency



One architecture, one backbone, one thread count — no capacity difference.

Weights 36.2 MiB against a 2 MiB shared L3: never resident, streamed from DRAM every inference.