


Memory traffic, not arithmetic, sets the latency



Single model, single backbone, single thread count — no capacity confound.

Weights 36.2 MiB against a 2 MiB shared L3: never resident, streamed from DRAM every inference.